

An industry, multi-system, PROV-DM-based architecture for Explainable AI: Application Track

Nicholas J. Car
*SURROUND Pty Ltd &
Australian National University*
nicholas.car@surroundaustralia.com

Robert A. Atkinson
*SURROUND Pty Ltd &
Open Geospatial Consortium*
rob.atkinson@surroundaustralia.com

Abstract

SURROUND Australia Pty Ltd is a small technology company focused on explainable AI and knowledge management products. At the core of our company operations is standardised provenance tracking using the PROV Data Model (PROV-DM). Using standardised provenance throughout our company lets us both assure customers that any results we produce for them are explainable, regardless of the specific systems we employ, and also allows us to easily learn about multi-system performance.

We generate PROV-DM-based provenance in all our important business workflows by using either our in-house *ProvWF* tool or workflows within our proprietary *SURROUND Ontology Platform* (SOP) tool. We also implement PROV-DM within our project and company data which are mostly Knowledge Graph (KG) products. We “hand off” certain forms of provenance tracking - for some data and software - to Git-based systems but do ensure that they integrate with our PROV-DM model.

1 Introduction

SURROUND Australia Pty Ltd (“SURROUND”) is a small technology company founded 6 years ago. While aiming to supply mainstream AI and knowledge management products to government and private sector markets, SURROUND differentiates itself from competitors through sophisticated use of Semantic Web data due to the belief that it is the form that best preserves meaning over time, system & organisational change. Specific value propositions for SURROUND’s customers, based on Semantic Web data use, are:

- the expressivity of RDFS / OWL2 [1, 6] allows for infinitely complex, yet cohesive, data modelling
- the ability to directly reuse many existing, sophisticated, published models (ontologies)

- the extensibility of RDF graph-based data structures - no need to alter system schema as they grow
- the system-independence of Semantic Web data formats allowing for seamless under-the-hood system changes
- the ability to use a Semantic Web layer to act as a bridge between internal, siloed applications
- the ability to share data across organisational boundaries with no inter-organisational special data contracts (due to Semantic modelling of all data elements)
- the data validation power of modern constraints languages such as SHACL [2]
- the advanced reasoning capabilities of OWL & SHACL to infer new knowledge

Emergent from some of these points is SURROUND’s ability to provide comensurate provenance information across all our different systems.

In this paper we make no new research claim - this is an *Applications Track* paper - but we do claim to show “innovative use of provenance” and “the deployment of provenance-based solutions” that indicate a certain maturity of approach to the use of provenance for operational tasks. We expect this will be of interests to our industry peers and to academics wishing to know where industry implementers are up to, in order to consider next research steps.

We will overview our specific company-wide provenance systems, discuss two projects that use them, indicate why we’ve chosen certain PROV-related implementations over others and where we think some of the provenance standards need enhancement for our purposes.

2 Simple provenance theory, complex practice

In general, extraction of useful information present within heterogeneous or large-scale data contexts may be performed in several ways. If some of the data has some known structure

then queries can be used to select relevant subsets. The trivial form of this is searching against text content, with various degrees of sophistication. These may involve statistical techniques to identify patterns in the data. SURROUND uses Machine Learning (ML) approaches to train systems to correlate or discover information based on various patterns. We also use Semantic or Knowledge Graph (KG)-based contextual information to improve performance. Conversely, we also use ML approaches to infer structure. Some of our projects include Human-in-the-loop (HITL) activities too, to refine, record and improve training of systems. Performing these tasks requires us to implement complex, hybrid systems that use reasoning and ML to create ways to to organise and retrieve information from complex projects, two of which are summarised in the next sections.

To record provenance systematically across our multiple systems and our various data processing domains requires a sensible provenance reference model and detailed system/data translation/communication/integration technical work. Since the advent and then widespread adoption of the PROV Data Model (PROV-DM) [4] and it's Semantic Web form PROV-O [3], we have a provenance reference model that SURROUND applies "everywhere" (as far as we practically can). The model is flexible enough for our use, with some small extensions, within our many systems and for many scenarios and cohesive enough for systematic use.

The technical implementation challenges we face are characterised as:

1. **Shared entity identification** - between different systems as they work on and pass the entities
2. **Granularity** - having useful levels of knowledge detail whilst retaining ability to have overviews from the process or system level

Shared entity identification is well handled using the RDF meta model, which uses unique URIs for things (entities and others) and the Open World Assumption, as multiple systems can represent knowledge in RDF data structures and then join knowledge by referencing shared URIs. SURROUND not only communicates our provenance information in RDF but, for most projects, our primary data also. Since our project data, perhaps electronic records described using RDF metadata, and our provenance information describing that data's generation are both represented in RDF, we can identify entities across both information holdings too, not just across systems within each holding. For non-RDF information, such as Git-based software and data version information, we ensure that URIs are also used.

Architecturally, our project data and provenance can be integrated across multiple subsystems if:

1. Object identity is established in KGs and they are accessed via APIs

2. Object identity is managed within the KG whenever "human in the loop" interactions are required
3. Processing subsystems preserve and report canonical object identities
4. Coherent sets of objects may be managed in specialised persistence systems, such as Git, as long as the description of datasets containing them are managed within KGs (see dataset granularity)
5. All processing reports' provenance, along with outputs, use our canonical model - PROV-DM
6. Processing elements are identified coherently in KGs

To achieve the above points, we have implemented systems and methods described in the next section. To demonstrate what such implementation allows for, we include Figure 1 which shows the user interface of the *SURROUND Ontology Platform* (SOP) displaying provenance information within a sankey diagram. The provenance information was generated according to PROV-DM/PROV-O by SURROUND's processing workflow tool *ProvWF* performing Named Entity Recognition against electronic records matching entities against some of SURROUND's Knowledge Graph products. In addition to displaying the provenance information in particular ways, SOP also managed it in bundles as *Managed Graphs* which, from SOP's point of view are yet another semantic asset for which provenance (and ownership, access etc.) is automatically stored.

The provenance granularity issue identified above can be summarised by examining a range of different types of processing that may typically occur in a heterogeneous systems and do occur within our projects. Table 1 lists processing functions, examples of them and (our) required granularity.

3 Company-wide provenance architecture

Figure 2 links types of IT project assets we work with to the tools we implement to record PROV-O-compliant provenance. Our major provenance tools and the actions they perform are:

- **SURROUND Ontology Platform (SOP)**¹
 - an enterprise data management system based on semantic data that is based on Top Quadrant's *EDG*²
 - SOP extends EDG adding management of semantic asset state and collections of them
 - SOP records PROV-DM provenance for all semantic asset actions

¹<https://surroundaustralia.com/sop>

²<https://www.topquadrant.com/products/topbraid-enterprise-data-governance/>

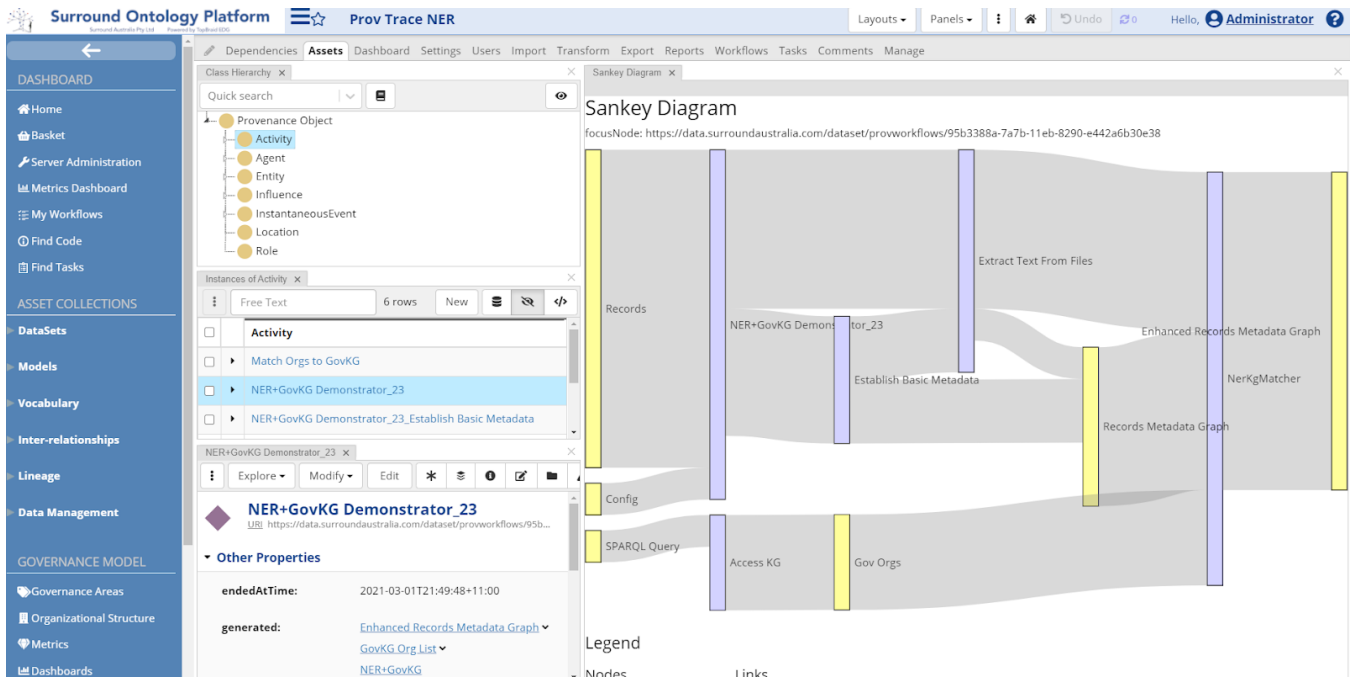


Figure 1: An example of a provenance trace from a processing workflow that uses elements of a knowledge graph, performs processing in cloud-hosted scalable services, generates augmented views of an input stream (performing Named Entity Recognition on a document set and annotating with elements from the knowledge graph), persists the results in the knowledge graph and integrates the provenance trace with the provenance trace generated by knowledge graph management.

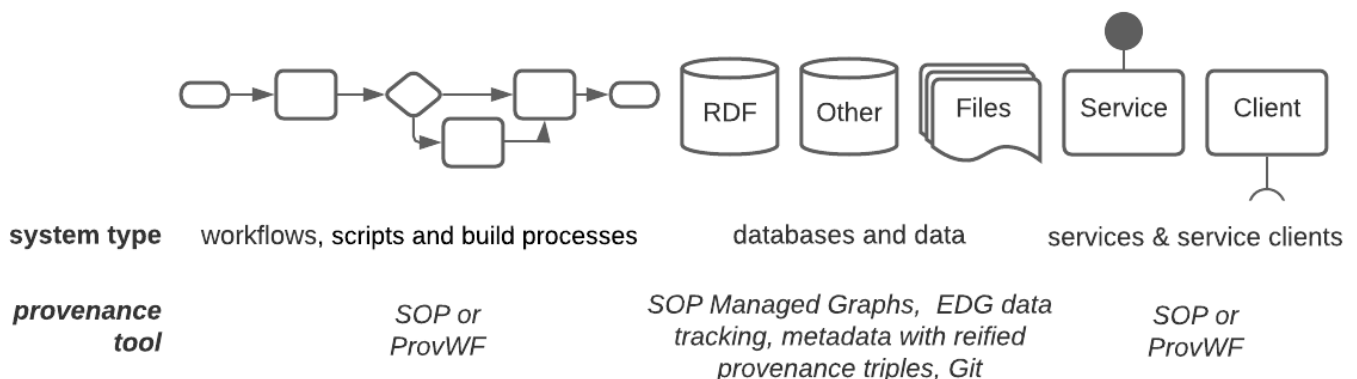


Figure 2: SURROUND's provenance tools linked to system type

Function	Examples	Granularity
Human-in-the-loop classification	ML Establishment of defs, Registration entities, Annotation, Classification for training	Statement, Reified statements
Database management, Data transformation	Making data instances sets available in a useful form	Dataset (table, spreadsheet, graph etc)
Query	Extraction of data subsets	Dataset, Resultset
Governance	Selecting particular datasets for use	Dataset
Bulk object processing	Indexing, classification, clustering	Whole-of-workflow
Document analysis	Making information elements in a document available to finer grained processes	Document, derived dataset
KG Management	Est'ment of state of complex, modular KGs, change tracking, support for automated updates	Graph (Dataset)

Table 1: A list of project functions, examples of them and (our) required provenance granularity

- **ProvWorkflow** (ProvWF)³
 - Python framework for creation of workflows
 - records PROV-DM provenance for actions performed by the workflow and data consumed/produced
 - supported by SURROUND’s *Block Library*
 - integrates with SOP by provenance bundle transfer, see Figure 3
- **Block Library**
 - SURROUND’s catalogue of ProvWF *Blocks*: PROV-DM Activity class objects
 - library stores many reusable functions within *Blocks*, such as KG API requests, NLP text processing etc.
- **Git**⁴
 - to track the versions of many assets - code, data etc. - in both public and private repositories
 - we have not yet found it necessary to materialise Git-to-PROV mappings, e.g. Git2PROV [5], but instead record URIs for entities managed in Git and refer to them in PROV-O data

In addition to these tools, we implement provenance tracking with general-purpose tools too, such as:

- **RDFlib**⁵
 - general-purpose Python RDF manipulation library

³<https://surroundaustralia.com/provWF>

⁴<https://git-scm.com/>

⁵<https://github.com/RDFLib/rdfliib>

- many of our data objects are RDF graphs
- we maintain RDFlib code blocks, e.g. to create reified provenance for RDF statements

4 Aspects of our provenance modelling

Much of our provenance modelling will be familiar to PROV users: chains of *Activities* and *Entities* associated with *Agents*. We do have several project scenarios which cause us to implement slight specialisations and two are given in project-specific case studies next.

4.1 Electronic Records assessment

A recent project of ours has applied Natural Language Processing and KGs to electronic Records classification for archiving: we have extracted elements of Records’ content, compared them with managed sources of *context* presented as KGs and used ML to learn how best to classify the Records.

Important in this project’s provenance capture has been the use of data services - our classification workflow system querying our own KG services - and the tracking of both whole-workflow and individual data statement provenance. Figure 4 shows our model for data service query tracking within a workflow. Whole-of-workflow provenance is needed to track what configurations produce what results. Individual data statement provenance is needed for classifications within a Record’s metadata in RDF so individual Record results can be verified. Workflow, metadata and metadata provenance are all stored in an RDF database and are crosslinked.

4.2 Semantic document querying

Another project of ours involved decomposing a large industry specification into not only structural but also semantic

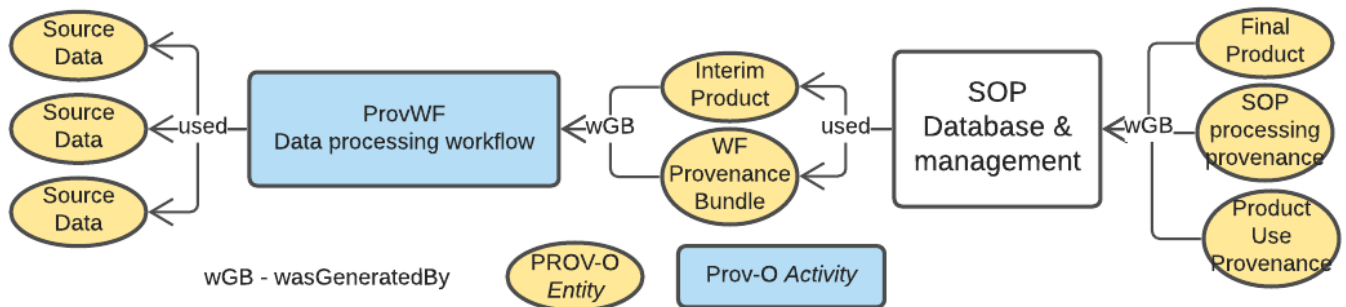


Figure 3: *ProvWF* is often used to generate RDF data - here the “Interim Product” - which can be supplied to *SOP* with an accompanying provenance Bundle. *SOP*, in turn, generates both Bundles of provenance for any actions on data it performs and also records usage provenance for products

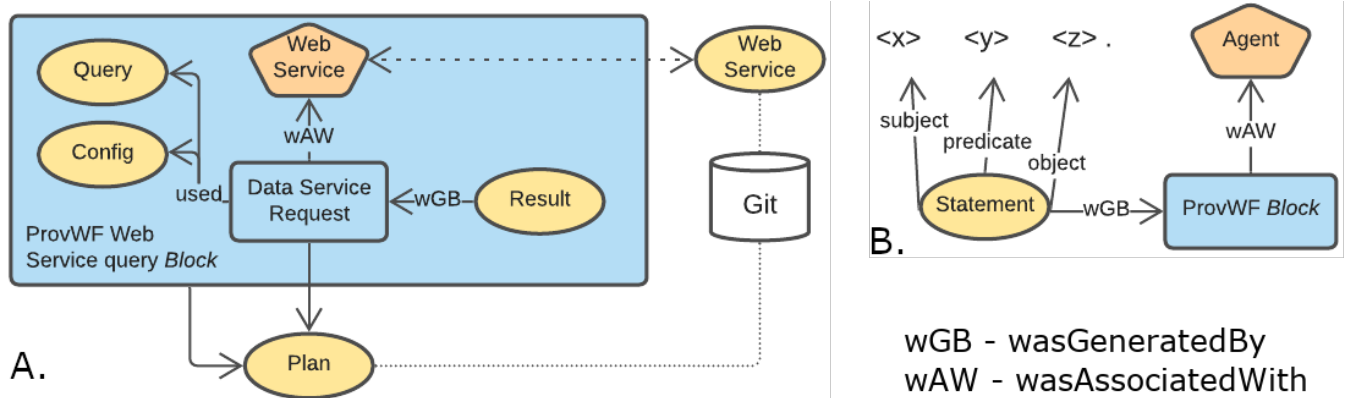


Figure 4: **A.** A *Service query block* from our *Block Library*, implemented within our *ProvWF* framework using a *Query* and other configuration (*Config*) to query a *Web Service* agent for a *Result*. Provenance for the web service itself, now considered an entity, is recorded in *Git* systems and referenced by each *ProvWF* execution. **B.** Reified provenance for a single RDF triple associated with the *ProvWF* Block instance that generated it.

elements (phrase meanings, synonyms, diagram descriptions, terminology lists, algorithm elements) to facilitate natural language and naïve searching of it.

5 Conclusions

SURROUND uses fairly standard patterns with PROV-DM to model provenance across multiple systems and within different IT domains. We have easily adapted, and slightly specialised, the model to provide provenance that gives our customers confidence in the results they receive from us, particularly important for complex AI/ML applications, but have had to put considerable effort into establishing technical interfaces, interactions and integrations between our tools. We have developed both dedicated provenance tools and provenance capability within generic tools.

References

- [1] Dan Brickley and R.V. Guha. RDF Schema 1.1. W3C Recommendation, World Wide Web Consortium, February 2014.
- [2] Holger Knublauch and Dimitris Kontokostas. Shapes Constraint Language (SHACL). W3C Recommendation, W3C RDF Data Shapes Working Group, 2017.
- [3] Timothy Lebo, Satya Sahoo, and Deborah McGuinness. PROV-O: The PROV Ontology. W3C Recommendation, W3C Provenance Working Group, 2013.
- [4] Luc Moreau and Paolo Missier. PROV-DM: The PROV Data Model. W3C Recommendation, World Wide Web Consortium, 2013.
- [5] Tom De Nies, Sara Magliacane, Ruben Verborgh, Sam Coppens, Paul Groth, Erik Mannens, and Rik Van de Walle. Git2PROV: Exposing Version Control System Content as W3C PROV. In *Proceedings of the 12th International Semantic Web Conference*, volume II. Springer.
- [6] W3C OWL Working Group. OWL 2 Web Ontology Language Document Overview (Second Edition). W3C Recommendation, 2012.